


Data Visualization in Python: Advanced Functionality in Seaborn

13 NOVEMBER 2015



Slater Stich is an Insight Data Labs Mentor and was previously a Staff Data Scientist at Square. He is currently pursuing a graduate degree in business at Stanford.

Seaborn is a Python data visualization library with an emphasis on statistical plots. The library is an excellent resource for common regression and distribution plots, but where Seaborn really shines is in its ability to visualize many different features at once.

In this post, we'll cover three of Seaborn's most useful functions:

`factorplot`, `pairplot`, and `jointgrid`. Going a step further, we'll show how we can get even more mileage out of these functions by stepping up to their even-more-powerful forms: `FacetGrid`, `PairGrid`, and `JointGrid`.

The Data

To showcase Seaborn, we'll use the [UCI "Auto MPG" data set](#). We did a bit of preprocessing: to see the details check out the [IPython notebook](#) that accompanies this post.

`factorplot` and `FacetGrid`

One of the most powerful features of Seaborn is the ability to easily build conditional plots; this let's us see what the data look like when segmented by one or more variables. The easiest way to do this is thorough `factorplot`. Let's say that we we're interested in how cars' MPG has varied over time. Not only can we easily see this in aggregate:

```
sns.factorplot(data=df, x="model_year", y="mpg")
```



But we can also segment by, say, region of origin:

```
sns.factorplot(data=df, x="model_year", y="mpg", col="origin")
```



What's so great `factorplot` is that rather than having to segment the data ourselves and make the conditional plots individually, Seaborn provides a convenient API for doing it all at once.

The `FacetGrid` object is a slightly more complex, but also more powerful, take on the same idea. Let's say that we wanted to see KDE plots of the MPG distributions, separated by country of origin:

```
g = sns.FacetGrid(df, col="origin")
g.map(sns.distplot, "mpg")
```



Or let's say that we wanted to see scatter plots of MPG against horsepower with the same origin segmentation:

```
g = sns.FacetGrid(df, col="origin")
g.map(plt.scatter, "horsepower", "mpg")
```




Using `FacetGrid`, we can map any plotting function onto each segment of our data. For example, above we gave `plt.scatter` to `g.map`, which tells Seaborn to apply the matplotlib `plt.scatter` function to each of segments in our data. We don't need to use `plt.scatter`, though; we can use any function that understands the input data. For example, we could draw regression plots instead:

```
g = sns.FacetGrid(df, col="origin")
g.map(sns.regplot, "horsepower", "mpg")
plt.xlim(0, 250)
plt.ylim(0, 60)
```



We can even segment by multiple variables at once, spreading some along the rows and some along the columns. This is very useful for producing comparing conditional distributions across interacting segmentations:

```
df['tons'] = (df.weight/2000).astype(int)
g = sns.FacetGrid(df, col="origin", row="tons")
g.map(sns.kdeplot, "horsepower", "mpg")
plt.xlim(0, 250)
plt.ylim(0, 60)
```



pairplot and PairGrid

While `factorplot` and `FacetGrid` are for drawing conditional plots of segmented data, `pairplot` and `PairGrid` are for showing the interactions between variables. For our car data set, we know that MPG, horsepower, and weight are probably going to be related; we also know that both these variable values *and* their relationships with one another, might vary by country of origin. Let's visualize all of that at once:

```
g = sns.pairplot(df[["mpg", "horsepower", "weight", "origin"]],
hue="origin", diag_kind="hist")

for ax in g.axes.flat:
    plt.setp(ax.get_xticklabels(), rotation=45)
```



As **FacetGrid** was a fuller version of **factorplot**, so **PairGrid** gives a bit more freedom on the same idea as **pairplot** by letting you control the individual plot types separately. Let's say, for example, that we're building regression plots, and we'd like to see both the original data and the residuals at once. **PairGrid** makes it easy:

```
g = sns.PairGrid(df[["mpg", "horsepower", "weight", "origin"]],
hue="origin")

g.map_upper(sns.regplot)
g.map_lower(sns.residplot)
g.map_diag(plt.hist)

for ax in g.axes.flat:
    plt.setp(ax.get_xticklabels(), rotation=45)

g.add_legend()
g.set(alpha=0.5)
```



We were able to control three regions (the diagonal, the lower-left triangle, and the upper-right triangle) separately. Again, you can pipe in any plotting function that understands the data it's given.

jointplot and JointGrid

The final Seaborn objects we'll talk about are `jointplot` and `JointGrid`; these features let you easily view both a joint distribution and its marginals at once. Let's say, for example, that aside from being interested in how MPG and horsepower are distributed individually, we're also interested in their joint distribution:

```
sns.jointplot("mpg", "horsepower", data=df, kind='kde')
```



As before, `JointGrid` gives you a bit more control by letting you map the marginal and joint data separately. For example:

```
g = sns.JointGrid(x="horsepower", y="mpg", data=df)
g.plot_joint(sns.regplot, order=2)
g.plot_marginals(sns.distplot)
```

|

Bottom Line

Seaborn is a great Python visualization library, and some of its most powerful features are:

- `factorplot` and `FacetGrid`,
- `pairplot` and `PairGrid`,

- [jointplot](#) and [JointGrid](#)

If you're new to Seaborn, the [official Seaborn tutorial](#) is a great place to start learning about simpler, but also extremely useful, functions such as

[distplot](#), [regplot](#), and the other component functions we used above.


For more data visualization tutorials, [join our mailing list](#).

Want to learn more advanced data visualization topics?

Join Insight for a two-day advanced [Data Visualization Lab](#) led by tech industry experts.

Interested in transitioning to a career in data science?

Learn more about the [Insight Data Science Fellows Program](#).



Slater Stich

Read [more posts](#) by this author.

Share this post



Be the first to get new Data Labs tutorials.

Your email address

SUBSCRIBE

Insight

DATA SCIENCE FELLOWSHIP

DATA ENGINEERING FELLOWSHIP

HEALTH DATA FELLOWSHIP

Data Labs

VISUALIZATION LAB

SPARK LAB

BLOG

©2015 Insight Data Labs



datalabs@insightdatascience.com